

RemoveNodeMAB-Different_bandit

March 22, 2017

Compute the performance of MAB methods

```
In [1]: import numpy as np
import time
import sys
import matplotlib.pyplot as plt
from sklearn import metrics
%matplotlib inline
plt.rcParams['figure.figsize'] = (15, 6)
```

0.1 Load BOKEH libariry

```
In [2]: from bokeh.layouts import row, gridplot
from bokeh.plotting import figure, output_notebook, show
from bokeh.models import Legend
```

```
#####
TOOLS = 'box_zoom,box_select,crosshair,resize,reset,lasso_select,pan,save,poly_select,ta
output_notebook()
#####
```

1 Load the data

```
In [3]: X_train = np.load('X_train.npy')
y_train = np.load('y_train.npy')
X_test = np.load('X_test.npy')
y_test = np.load('y_test.npy')
X_deploy = np.load('X_deploy.npy')
y_deploy = np.load('y_deploy.npy')

print('Number of training examples',len(X_train))
print('Number of validation examples',len(X_test))
print('Number of testing examples',len(X_deploy))
```

```
Number of training examples 2944
Number of validation examples 736
Number of testing examples 921
```

```
In [4]: exec(open("core.py").read()) # python 3x
        #exec(compile(open('core.py', "rb").read(), 'core.py', 'exec'))

        #execfile("/Users/saleameen/Desktop/banditsbook/python/core.py") # python 2.7
```


1.1 Run Thompson Sampling pruning Algorithm

```
In [5]: algo = Thompson_Sampling([], [])
        Alg_name = 'Thompson_Sampling Algorithm'
        path = './Thompson_Sampling/'
        sys.path.append("./Thompson_Sampling")
        exec(open("mnist_cnnFORTESTING.py").read())
```

Using Theano backend.

```
736 test samples
Test score: 0.338247084349
Test accuracy: 0.944625407166
The time for running this method is 1.7069087028503418 seconds
Finsh playing start pruning:
Test accuracy after pruning: 0.944625407166
Test accuracy after pruning: 0.944625407166
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.940282301846
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.941368078176
Test accuracy after pruning: 0.941368077593
Test accuracy after pruning: 0.941368077593
Test accuracy after pruning: 0.941368077593
Test accuracy after pruning: 0.942453853924
Test accuracy after pruning: 0.943539630254
Test accuracy after pruning: 0.940282302299
Test accuracy after pruning: 0.941368078629
Test accuracy after pruning: 0.945711183949
Test accuracy after pruning: 0.944625407619
Test accuracy after pruning: 0.921824104688
Test accuracy after pruning: 0.921824104688
Test accuracy after pruning: 0.922909880435
Test accuracy after pruning: 0.921824104105
Test accuracy after pruning: 0.918566775115
Test accuracy after pruning: 0.912052117134
```

Test accuracy after pruning: 0.895765472636
 Test accuracy after pruning: 0.893593919394
 Test accuracy after pruning: 0.889250814073
 Test accuracy after pruning: 0.890336590403
 Test accuracy after pruning: 0.891422366734
 Test accuracy after pruning: 0.885993486119
 Test accuracy after pruning: 0.834961998605
 Test accuracy after pruning: 0.766558089228
 Test accuracy after pruning: 0.753528772684
 Test accuracy after pruning: 0.375678610061
 Test accuracy after pruning: 0.375678610061
 Test accuracy after pruning: 0.375678610061
 Test accuracy after pruning: 0.375678610061



1.2 Run UCB1 pruning Algorithm

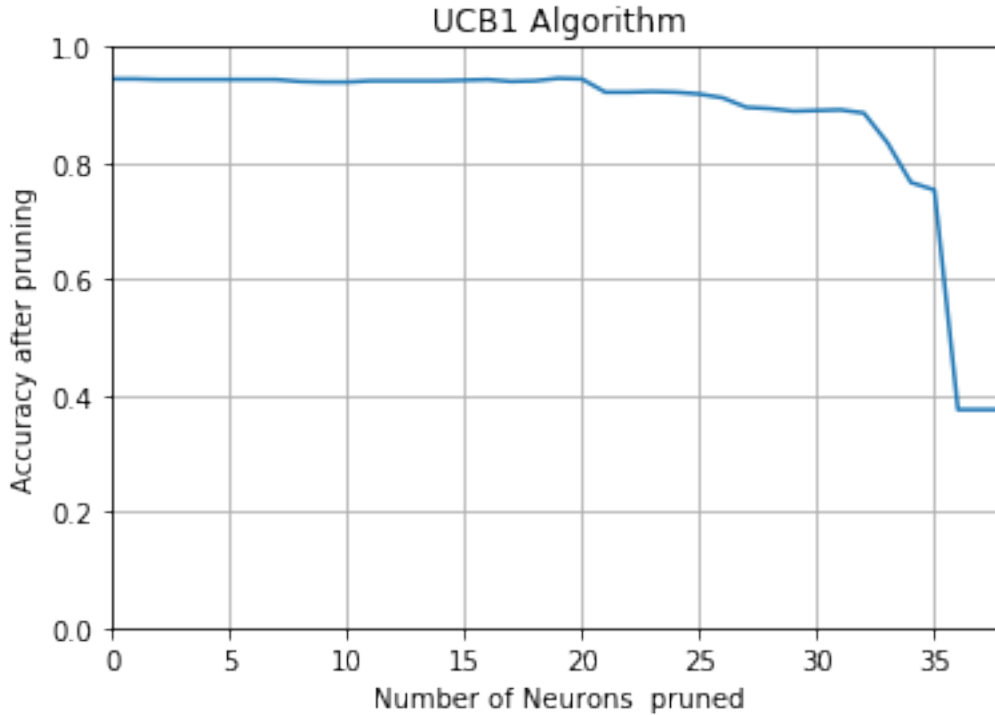
```

In [6]: algo = UCB1([], [])
        Alg_name = 'UCB1 Algorithm'
        path = './UCB1/'
        sys.path.append("./UCB1")
        exec(open("mnist_cnnFORTESTING.py").read())
  
```

736 test samples

Test score: 0.338247084349

Test accuracy: 0.944625407166
The time for running this method is 1.6735870838165283 seconds
Finsh playing start pruning:
Test accuracy after pruning: 0.944625407166
Test accuracy after pruning: 0.944625407166
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.940282301846
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.941368078176
Test accuracy after pruning: 0.941368077593
Test accuracy after pruning: 0.941368077593
Test accuracy after pruning: 0.941368077593
Test accuracy after pruning: 0.942453853924
Test accuracy after pruning: 0.943539630254
Test accuracy after pruning: 0.940282302299
Test accuracy after pruning: 0.941368078629
Test accuracy after pruning: 0.945711183949
Test accuracy after pruning: 0.944625407619
Test accuracy after pruning: 0.921824104688
Test accuracy after pruning: 0.921824104688
Test accuracy after pruning: 0.922909880435
Test accuracy after pruning: 0.921824104105
Test accuracy after pruning: 0.918566775115
Test accuracy after pruning: 0.912052117134
Test accuracy after pruning: 0.895765472636
Test accuracy after pruning: 0.893593919394
Test accuracy after pruning: 0.889250814073
Test accuracy after pruning: 0.890336590403
Test accuracy after pruning: 0.891422366734
Test accuracy after pruning: 0.885993486119
Test accuracy after pruning: 0.834961998605
Test accuracy after pruning: 0.766558089228
Test accuracy after pruning: 0.753528772684
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061



2 Run Annealing Epsilon Greedy pruning Algorithm

```
In [7]: algo = AnnealingEpsilonGreedy([], [])
        Alg_name = 'Annealing Epsilon Greedy Algorithm'
        path = './AnnealingEpsilonGreedy/'
        sys.path.append("./AnnealingEpsilonGreedy")
        exec(open("mnist_cnnFORTESTING.py").read())
```

736 test samples

Test score: 0.338247084349

Test accuracy: 0.944625407166

The time for running this method is 1.6822741031646729 seconds

Finsh playing start pruning:

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.939196524933

Test accuracy after pruning: 0.940282301263

Test accuracy after pruning: 0.941368078176

Test accuracy after pruning: 0.942453854506

Test accuracy after pruning: 0.941368078176

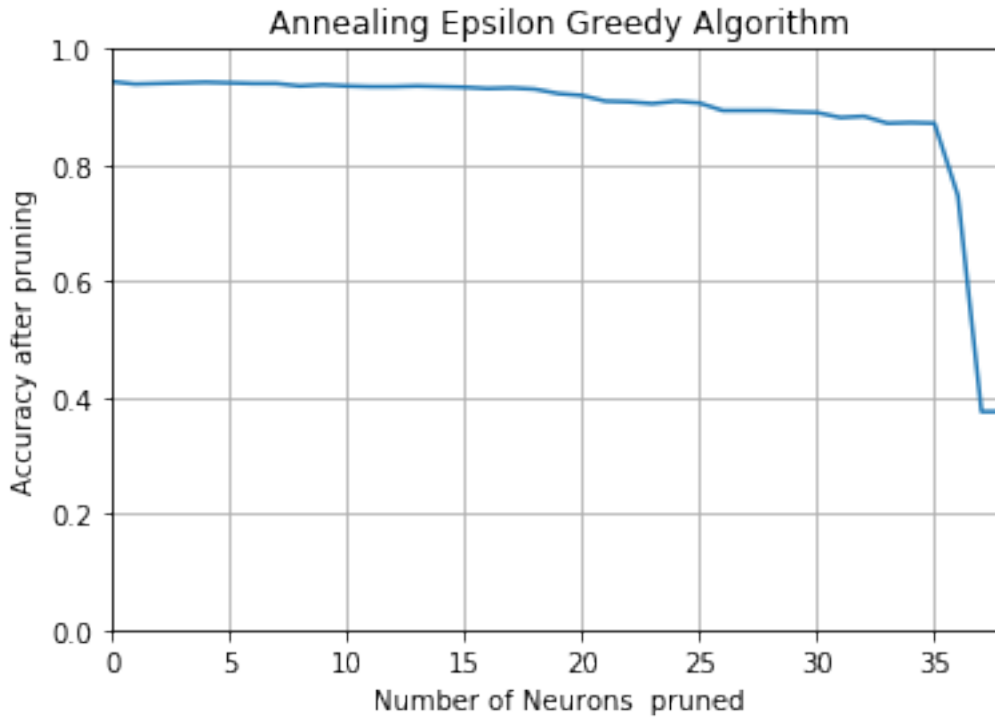
Test accuracy after pruning: 0.940282301846

Test accuracy after pruning: 0.940282301846

Test accuracy after pruning: 0.935939196526

Test accuracy after pruning: 0.938110749186

Test accuracy after pruning: 0.935939196526
Test accuracy after pruning: 0.934853420195
Test accuracy after pruning: 0.934853420195
Test accuracy after pruning: 0.935939196526
Test accuracy after pruning: 0.934853419613
Test accuracy after pruning: 0.933767643283
Test accuracy after pruning: 0.931596091658
Test accuracy after pruning: 0.932681867988
Test accuracy after pruning: 0.930510314293
Test accuracy after pruning: 0.922909879853
Test accuracy after pruning: 0.919652550862
Test accuracy after pruning: 0.909880564927
Test accuracy after pruning: 0.908794788015
Test accuracy after pruning: 0.905537459607
Test accuracy after pruning: 0.909880563892
Test accuracy after pruning: 0.906623234902
Test accuracy after pruning: 0.893593918941
Test accuracy after pruning: 0.893593919976
Test accuracy after pruning: 0.893593919976
Test accuracy after pruning: 0.891422367316
Test accuracy after pruning: 0.890336590403
Test accuracy after pruning: 0.881650379763
Test accuracy after pruning: 0.883821932423
Test accuracy after pruning: 0.871878393828
Test accuracy after pruning: 0.872964170158
Test accuracy after pruning: 0.871878392792
Test accuracy after pruning: 0.747014114704
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061



3 Run Epsilon Greedy pruning Algorithm

```
In [8]: epsilon = 0.9 # epsilon = (0,1)
        algo = EpsilonGreedy(epsilon, [], [])
        Alg_name = 'Epsilon Greedy Algorithm'
        path = './EpsilonGreedy/'
        sys.path.append("./AnnealingEpsilonGreedy")
        exec(open("mnist_cnnFORTESTING.py").read())
```

736 test samples

Test score: 0.338247084349

Test accuracy: 0.944625407166

The time for running this method is 1.5948808193206787 seconds

Finsh playing start pruning:

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

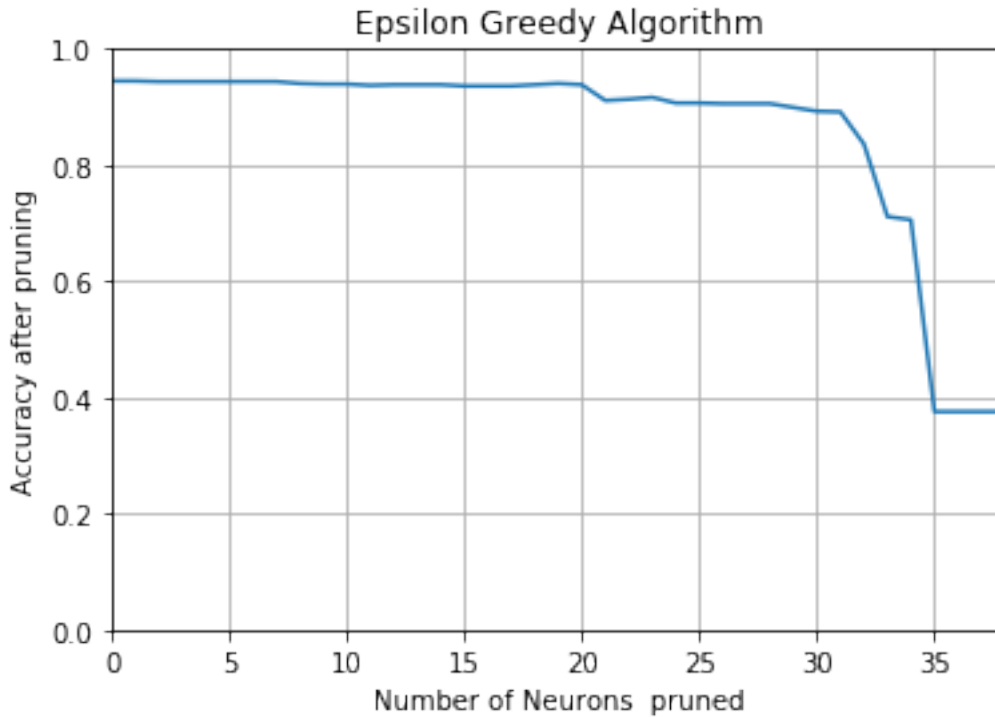
Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.940282301846

Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.937024972273
Test accuracy after pruning: 0.938110748603
Test accuracy after pruning: 0.938110748603
Test accuracy after pruning: 0.938110748603
Test accuracy after pruning: 0.935939196979
Test accuracy after pruning: 0.935939196396
Test accuracy after pruning: 0.935939196396
Test accuracy after pruning: 0.938110748474
Test accuracy after pruning: 0.940282301716
Test accuracy after pruning: 0.938110749056
Test accuracy after pruning: 0.910966340675
Test accuracy after pruning: 0.913137893335
Test accuracy after pruning: 0.916395222908
Test accuracy after pruning: 0.906623235355
Test accuracy after pruning: 0.906623235355
Test accuracy after pruning: 0.905537459025
Test accuracy after pruning: 0.905537459025
Test accuracy after pruning: 0.905537459025
Test accuracy after pruning: 0.899022801044
Test accuracy after pruning: 0.892508143064
Test accuracy after pruning: 0.891422366734
Test accuracy after pruning: 0.836047774935
Test accuracy after pruning: 0.711183496265
Test accuracy after pruning: 0.705754614614
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061



4 Run Exp3 pruning Algorithm

```
In [9]: exp3_gamma = 0.2 #exp3_gamma in [0.1, 0.2, 0.3, 0.4, 0.5]
        algo = Exp3(exp3_gamma, [])
        Alg_name = 'Exp3 Algorithm'
        path = './Exp3/'
        sys.path.append("./EpsilonGreedy")
        exec(open("mnist_cnnFORTESTING.py").read())
```

736 test samples

Test score: 0.338247084349

Test accuracy: 0.944625407166

The time for running this method is 1.6179580688476562 seconds

Finsh playing start pruning:

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

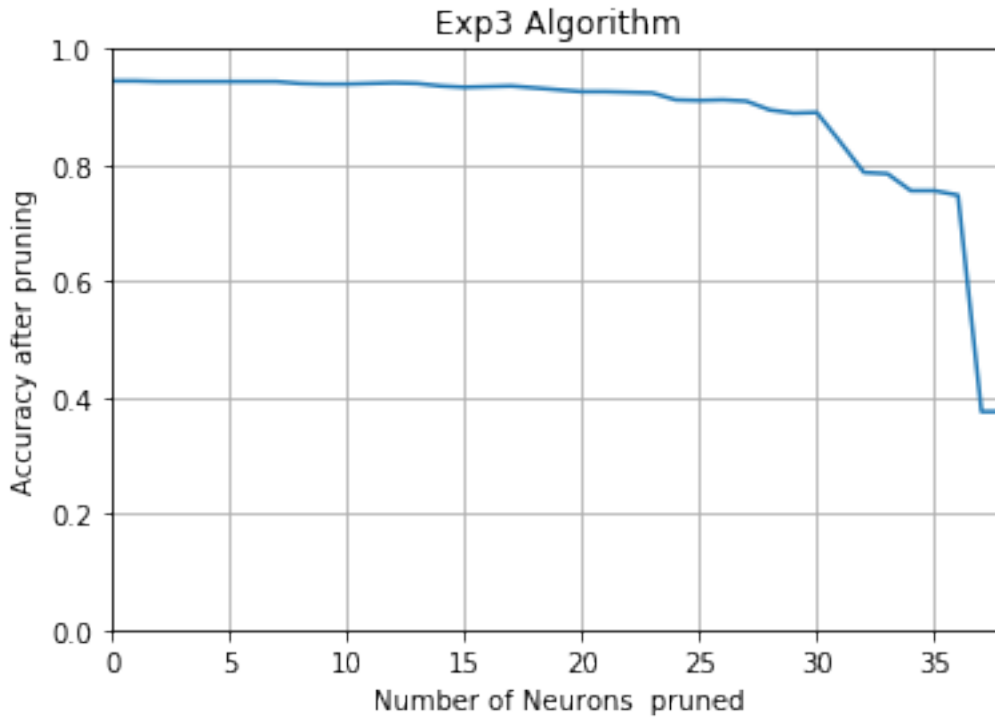
Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.940282301846

Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.940282301846
Test accuracy after pruning: 0.941368078176
Test accuracy after pruning: 0.940282301846
Test accuracy after pruning: 0.935939196526
Test accuracy after pruning: 0.933767643865
Test accuracy after pruning: 0.934853420195
Test accuracy after pruning: 0.935939196526
Test accuracy after pruning: 0.932681867535
Test accuracy after pruning: 0.929424537963
Test accuracy after pruning: 0.926167208972
Test accuracy after pruning: 0.926167208972
Test accuracy after pruning: 0.925081433678
Test accuracy after pruning: 0.923995657348
Test accuracy after pruning: 0.912052117134
Test accuracy after pruning: 0.910966340804
Test accuracy after pruning: 0.912052116552
Test accuracy after pruning: 0.909880563892
Test accuracy after pruning: 0.894679695724
Test accuracy after pruning: 0.889250814073
Test accuracy after pruning: 0.890336590403
Test accuracy after pruning: 0.83930510289
Test accuracy after pruning: 0.787187840082
Test accuracy after pruning: 0.785016287422
Test accuracy after pruning: 0.755700325927
Test accuracy after pruning: 0.755700325345
Test accuracy after pruning: 0.748099891617
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061



5 Run Softmax pruning Algorithm

```
In [10]: temperature = 0.9
         algo = Softmax(temperature, [], [])
         Alg_name = 'Softmax Algorithm'
         path = './Softmax/'
         sys.path.append("./Softmax")
         exec(open("mnist_cnnFORTESTING.py").read())
```

736 test samples

Test score: 0.338247084349

Test accuracy: 0.944625407166

The time for running this method is 1.5992538928985596 seconds

Finsh playing start pruning:

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

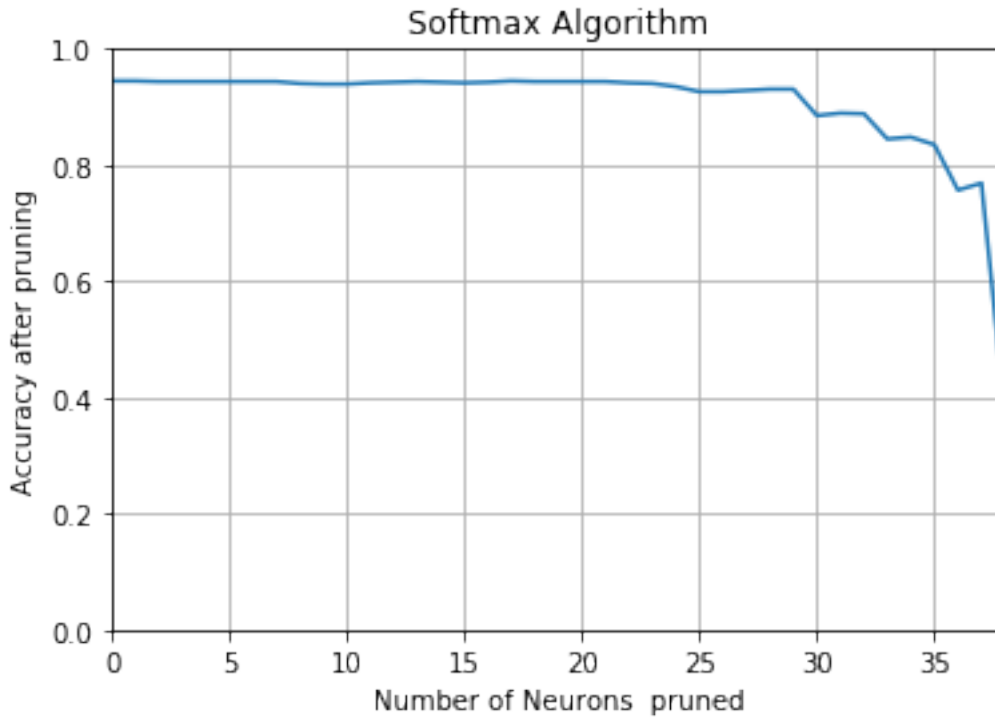
Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.940282301846

Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.941368078176
Test accuracy after pruning: 0.942453854506
Test accuracy after pruning: 0.943539630836
Test accuracy after pruning: 0.942453854506
Test accuracy after pruning: 0.941368078176
Test accuracy after pruning: 0.942453854506
Test accuracy after pruning: 0.944625406584
Test accuracy after pruning: 0.943539630254
Test accuracy after pruning: 0.943539630254
Test accuracy after pruning: 0.943539631289
Test accuracy after pruning: 0.943539631289
Test accuracy after pruning: 0.941368078629
Test accuracy after pruning: 0.940282302299
Test accuracy after pruning: 0.934853420648
Test accuracy after pruning: 0.926167209425
Test accuracy after pruning: 0.926167209425
Test accuracy after pruning: 0.928338762086
Test accuracy after pruning: 0.930510314746
Test accuracy after pruning: 0.930510314746
Test accuracy after pruning: 0.884907708753
Test accuracy after pruning: 0.889250814073
Test accuracy after pruning: 0.888165038779
Test accuracy after pruning: 0.844733985576
Test accuracy after pruning: 0.847991314566
Test accuracy after pruning: 0.834961998605
Test accuracy after pruning: 0.756786102257
Test accuracy after pruning: 0.768729641888
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061



6 Run Annealing Softmax pruning Algorithm

```
In [11]: algo = AnnealingSoftmax([], [])
         Alg_name = 'Annealing Softmax Algorithm'
         path = './AnnealingSoftmax/'
         sys.path.append("./AnnealingSoftmax")
         exec(open("mnist_cnnFORTESTING.py").read())
```

736 test samples

Test score: 0.338247084349

Test accuracy: 0.944625407166

The time for running this method is 1.604109764099121 seconds

Finsh playing start pruning:

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

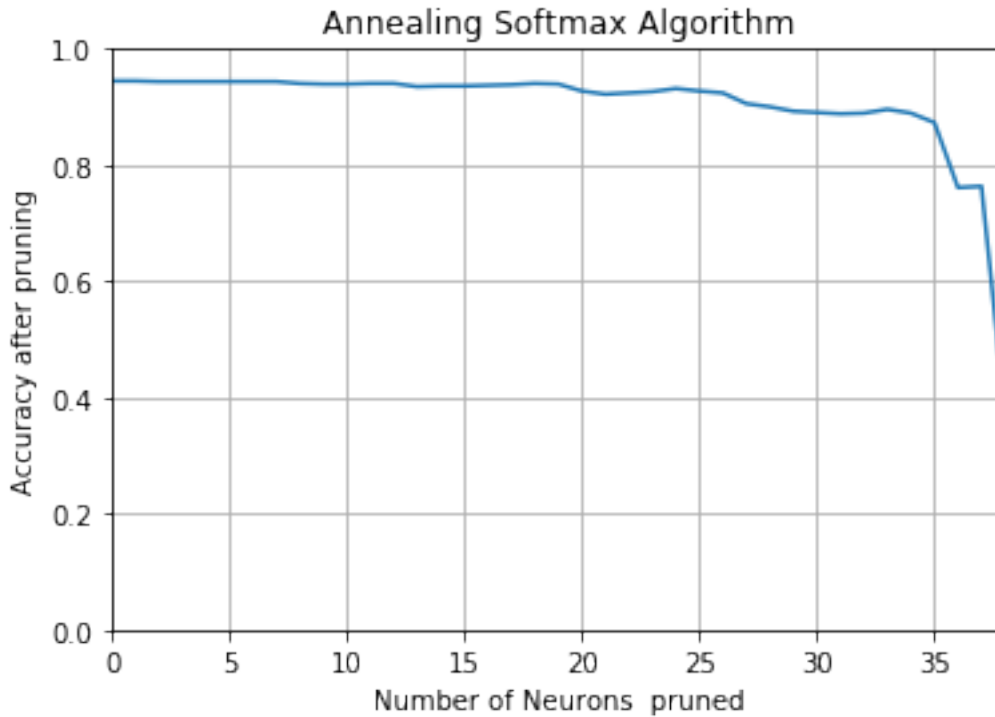
Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.940282301846

Test accuracy after pruning: 0.939196525516

Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.940282301846
Test accuracy after pruning: 0.940282301263
Test accuracy after pruning: 0.934853420648
Test accuracy after pruning: 0.935939196979
Test accuracy after pruning: 0.935939196979
Test accuracy after pruning: 0.937024973309
Test accuracy after pruning: 0.938110749639
Test accuracy after pruning: 0.940282301263
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.927252985173
Test accuracy after pruning: 0.921824104558
Test accuracy after pruning: 0.923995657218
Test accuracy after pruning: 0.926167209878
Test accuracy after pruning: 0.931596091076
Test accuracy after pruning: 0.927252985173
Test accuracy after pruning: 0.923995657218
Test accuracy after pruning: 0.905537459025
Test accuracy after pruning: 0.900108577374
Test accuracy after pruning: 0.892508143064
Test accuracy after pruning: 0.890336590403
Test accuracy after pruning: 0.888165038779
Test accuracy after pruning: 0.889250815109
Test accuracy after pruning: 0.895765473089
Test accuracy after pruning: 0.889250815109
Test accuracy after pruning: 0.872964169122
Test accuracy after pruning: 0.761129207577
Test accuracy after pruning: 0.763300760238
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061



7 Run Hedge pruning Algorithm

```
In [12]: eta = 0.9 # eta in [.5, .8, .9, 1, 2]
        algo = Hedge(eta, [], [])
        Alg_name = 'Hedge Algorithm'
        path = './Hedge/'
        sys.path.append("./Hedge")
        exec(open("mnist_cnnFORTESTING.py").read())
```

736 test samples

Test score: 0.338247084349

Test accuracy: 0.944625407166

The time for running this method is 1.6032681465148926 seconds

Finsh playing start pruning:

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.944625407166

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

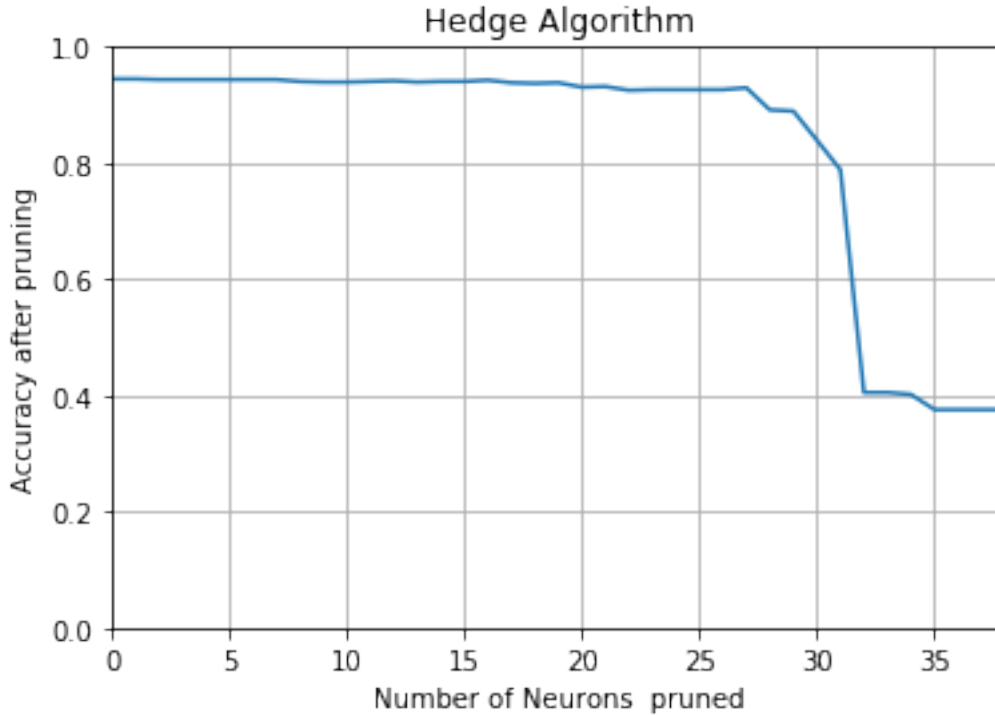
Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.943539630836

Test accuracy after pruning: 0.940282301846

Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.939196525516
Test accuracy after pruning: 0.940282301846
Test accuracy after pruning: 0.941368078176
Test accuracy after pruning: 0.939196524933
Test accuracy after pruning: 0.940282301263
Test accuracy after pruning: 0.940282301846
Test accuracy after pruning: 0.942453854506
Test accuracy after pruning: 0.938110749186
Test accuracy after pruning: 0.937024972856
Test accuracy after pruning: 0.938110748603
Test accuracy after pruning: 0.930510315328
Test accuracy after pruning: 0.931596091658
Test accuracy after pruning: 0.925081433095
Test accuracy after pruning: 0.926167209425
Test accuracy after pruning: 0.926167209425
Test accuracy after pruning: 0.926167208843
Test accuracy after pruning: 0.926167208843
Test accuracy after pruning: 0.929424537833
Test accuracy after pruning: 0.891422366734
Test accuracy after pruning: 0.889250814073
Test accuracy after pruning: 0.83930510289
Test accuracy after pruning: 0.788273616412
Test accuracy after pruning: 0.404994571151
Test accuracy after pruning: 0.404994571151
Test accuracy after pruning: 0.401737241983
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061
Test accuracy after pruning: 0.375678610061



8 Compare the accuracy of the models

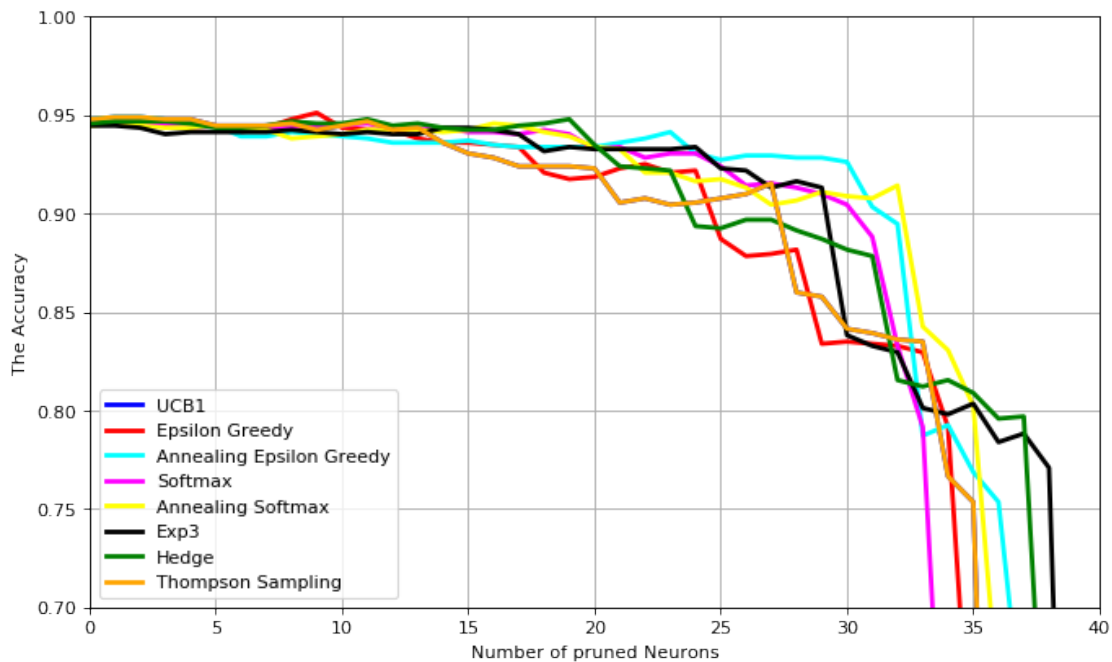
```
In [13]: ucb1 = np.load('/Users/saleem/Desktp/banditsbook/python/UCB1/AccuracyAfterPrune.n
EpsilonGreedy = np.load('/Users/saleem/Desktp/banditsbook/python/EpsilonGreedy/Acc
AnnealingEpsilonGreedy = np.load('/Users/saleem/Desktp/banditsbook/python/Annealin
Softmax = np.load('/Users/saleem/Desktp/banditsbook/python/Softmax/AccuracyAfterPr
AnnealingSoftmax = np.load('/Users/saleem/Desktp/banditsbook/python/AnnealingSoftm
Exp3 = np.load('/Users/saleem/Desktp/banditsbook/python/Exp3/AccuracyAfterPrune.n
Hedge = np.load('/Users/saleem/Desktp/banditsbook/python/Hedge/AccuracyAfterPrune
ThompsonSampling = np.load('/Users/saleem/Desktp/banditsbook/python/Thompson_Sampl
Accuracy = np.load('AccuracyBeforePruning.npy')
```

```
In [14]: fig = plt.figure(figsize=(10, 6), dpi=80)
ax = fig.add_subplot(111)
N = len(ucb1)
## necessary variables
ind = np.arange(N) # the x locations for the groups
plt.plot(ind , ucb1 , color="blue", linewidth=2.5, linestyle="-", label="UCB1")
plt.plot(ind , EpsilonGreedy, color="red", linewidth=2.5, linestyle="-", label="Epsilon
plt.plot(ind , AnnealingEpsilonGreedy, color="cyan", linewidth=2.5, linestyle="-", labe
plt.plot(ind , Softmax, color="magenta", linewidth=2.5, linestyle="-", label="Softmax")
plt.plot(ind , AnnealingSoftmax, color="yellow", linewidth=2.5, linestyle="-", label="A
plt.plot(ind , Exp3, color="black", linewidth=2.5, linestyle="-", label="Exp3")
```

```

plt.plot(ind , Hedge, color="green", linewidth=2.5, linestyle="-", label="Hedge")
plt.plot(ind , ThompsonSampling, color="orange", linewidth=2.5, linestyle="-", label="T")
plt.legend(loc = 3)
plt.axis([0, 40, 0.7, 1])
plt.xlabel('Number of pruned Neurons')
plt.ylabel('The Accuracy')
plt.grid(True)
plt.show()

```



```

In [15]: p1 = figure(title="The Performance over the number of neurons' pruned", tools=TOOLS)

#p1.circle(ind, ucb1, legend="ucb1", color="orange")
p1.line(ind, ucb1, legend="ucb1", line_color="orange", line_width=2)

#p1.square(ind, EpsilonGreedy, legend="Epsilon Greedy", fill_color=None, line_color="red")
p1.line(ind, EpsilonGreedy, legend="Epsilon Greedy", line_color="red", line_width=2)

#p1.ellipse(ind, AnnealingEpsilonGreedy, legend="Annealing Epsilon Greedy", line_color="blue")
p1.line(ind, AnnealingEpsilonGreedy, legend="Annealing Epsilon Greedy", line_color="blue", line_width=2)

#p1.diamond(ind, Softmax, legend="Softmax", line_color="green")
p1.line(ind, Softmax, legend="Softmax", line_color="green", line_width=2)

#p1.arc(ind, AnnealingSoftmax, legend="Annealing Softmax", line_color="grey", end_angle=90)
p1.line(ind, AnnealingSoftmax, legend="Annealing Softmax", line_color="grey", line_width=2)

```

```

#p1.oval(ind, Exp3, legend="Exp3", line_color="black", height=0.01, width=0.01)
p1.line(ind, Exp3, legend="Exp3", line_color="black", line_width=2)

#p1.arc(ind, Hedge, legend="Hedge", line_color="yellow")
#p1.triangle(ind, Hedge, legend="Hedge", line_color="yellow")
p1.line(ind, Hedge, legend="Hedge", line_color="yellow", line_width=2)

#p1.square_cross(ind, ThompsonSampling, legend="Thompson Sampling", line_color="pink")
p1.line(ind, ThompsonSampling, legend="Thompson Sampling", line_color="pink", line_widt

#p1.line(ind, Exp3, legend="2*sin(x)", line_dash=(4, 4), line_color="orange", line_widt
#p1.square(ind, Hedge, legend="3*sin(x)", fill_color=None, line_color="brown")
p1.title.align = "center"

show(p1)
#show(gridplot(p1, p2, ncols=2, plot_width=400, plot_height=400)) # open a browser

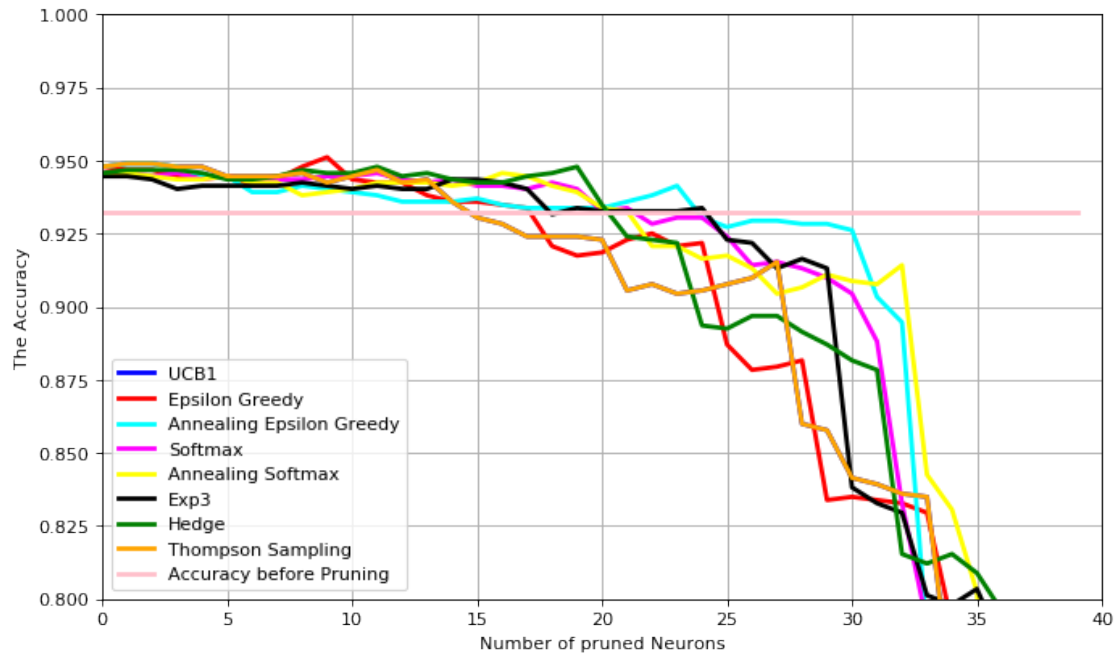
```

8.1 Comparing All algorithms with the model before pruning

```

In [16]: fig = plt.figure(figsize=(10, 6), dpi=80)
ax = fig.add_subplot(111)
N = len(ucb1)
Acc = [Accuracy for col in range(N)]
## necessary variables
ind = np.arange(N) # the x locations for the groups
plt.plot(ind, ucb1, color="blue", linewidth=2.5, linestyle="-", label="UCB1")
plt.plot(ind, EpsilonGreedy, color="red", linewidth=2.5, linestyle="-", label="Epsilon")
plt.plot(ind, AnnealingEpsilonGreedy, color="cyan", linewidth=2.5, linestyle="-", label="AnnealingEpsilonGreedy")
plt.plot(ind, Softmax, color="magenta", linewidth=2.5, linestyle="-", label="Softmax")
plt.plot(ind, AnnealingSoftmax, color="yellow", linewidth=2.5, linestyle="-", label="AnnealingSoftmax")
plt.plot(ind, Exp3, color="black", linewidth=2.5, linestyle="-", label="Exp3")
plt.plot(ind, Hedge, color="green", linewidth=2.5, linestyle="-", label="Hedge")
plt.plot(ind, ThompsonSampling, color="orange", linewidth=2.5, linestyle="-", label="ThompsonSampling")
plt.plot(ind, Acc, color="pink", linewidth=2.5, linestyle="-", label="Accuracy before pruning")
plt.legend(loc = 3)
plt.axis([0, 40, 0.8, 1])
plt.xlabel('Number of pruned Neurons')
plt.ylabel('The Accuracy')
plt.grid(True)
plt.show()

```



```
In [17]: p1 = figure(title="The Performance over the number of neurons' pruned", tools=TOOLS)

#p1.circle(ind, ucb1, legend="ucb1", color="orange")
p1.line(ind, ucb1, legend="ucb1", line_color="orange", line_width=2)

#p1.square(ind, EpsilonGreedy, legend="Epsilon Greedy", fill_color=None, line_color="red")
p1.line(ind, EpsilonGreedy, legend="Epsilon Greedy", line_color="red", line_width=2)

#p1.ellipse(ind, AnnealingEpsilonGreedy, legend="Annealing Epsilon Greedy", line_color="blue")
p1.line(ind, AnnealingEpsilonGreedy, legend="Annealing Epsilon Greedy", line_color="blue", line_width=2)

#p1.diamond(ind, Softmax, legend="Softmax", line_color="green")
p1.line(ind, Softmax, legend="Softmax", line_color="green", line_width=2)

#p1.arc(ind, AnnealingSoftmax, legend="Annealing Softmax", line_color="grey", end_angle=0.5)
p1.line(ind, AnnealingSoftmax, legend="Annealing Softmax", line_color="grey", line_width=2)

#p1.oval(ind, Exp3, legend="Exp3", line_color="black", height=0.01, width=0.01)
p1.line(ind, Exp3, legend="Exp3", line_color="black", line_width=2)

#p1.arc(ind, Hedge, legend="Hedge", line_color="yellow")
#p1.triangle(ind, Hedge, legend="Hedge", line_color="yellow")
p1.line(ind, Hedge, legend="Hedge", line_color="yellow", line_width=2)
```

```

#p1.square_cross(ind, ThompsonSampling, legend="Thompson Sampling", line_color="pink")
p1.line(ind, ThompsonSampling, legend="Thompson Sampling", line_color="pink", line_widt

p1.line(ind, Acc, legend="Accuracy", line_dash=(4, 4), line_color="orange", line_width=
#p1.square(ind, Hedge, legend="3*sin(x)", fill_color=None, line_color="brown")
p1.title.align = "center"

show(p1)
#show(gridplot(p1, p2, ncols=2, plot_width=400, plot_height=400)) # open a browser


```

8.2 UCB1

```

In [18]: fig = plt.figure(figsize=(10, 6), dpi=80)
ax = fig.add_subplot(111)
N = len(ucb1)
Acc = [Accuracy for col in range(N)]
## necessary variables
ind = np.arange(N) # the x locations for the groups
plt.plot(ind, ucb1, color="blue", linewidth=2.5, linestyle="-", label="UCB1")
plt.plot(ind, Acc, color="pink", linewidth=2.5, linestyle="-", label="Accuracy before
plt.legend(loc = 3)
plt.axis([0, 40, 0.7, 1])
plt.xlabel('Number of pruned Neurons')
plt.ylabel('The Accuracy')
plt.grid(True)
plt.show()

```



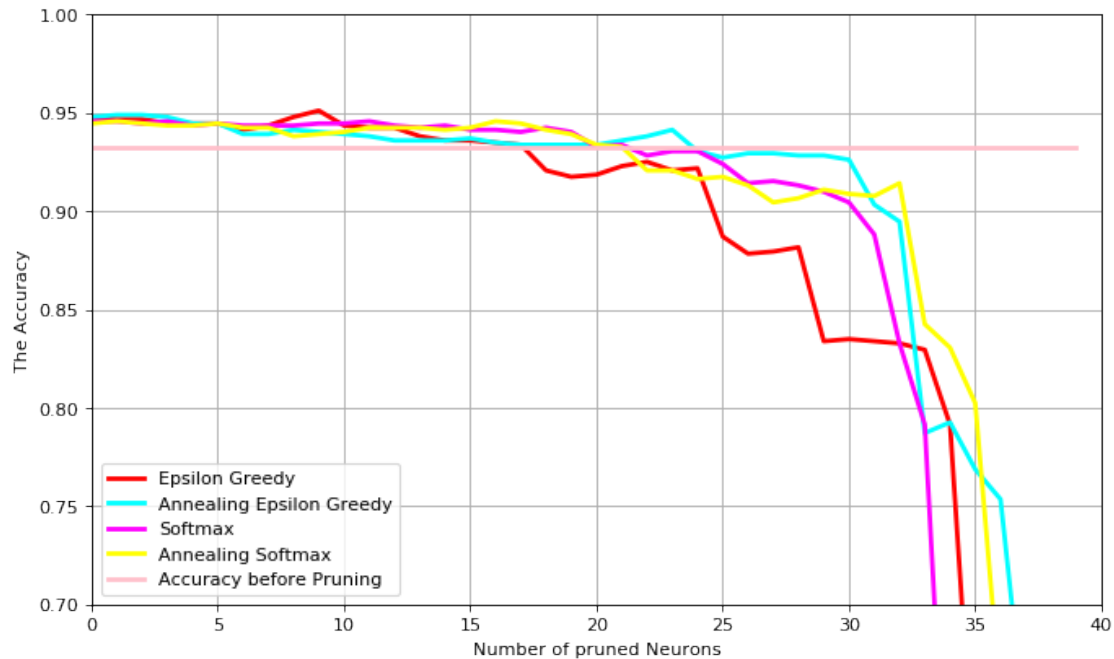
```
In [19]: p1 = figure(title="The Performance over the number of neurons' pruned", tools=TOOLS)

#p1.circle(ind, ucb1, legend="ucb1", color="orange")
p1.line(ind, ucb1, legend="ucb1", line_color="orange", line_width=2)
p1.line(ind, Acc, legend="Accuracy", line_dash=(4, 4), line_color="orange", line_width=2)
#p1.square(ind, Hedge, legend="3*sin(x)", fill_color=None, line_color="brown")
p1.title.align = "center"

show(p1)
#show(gridplot(p1, p2, ncols=2, plot_width=400, plot_height=400)) # open a browser
```

8.3 Epsilon greedy and Softmax

```
In [20]: fig = plt.figure(figsize=(10, 6), dpi=80)
ax = fig.add_subplot(111)
N = len(EpsilonGreedy)
Acc = [Accuracy for col in range(N)]
## necessary variables
ind = np.arange(N) # the x locations for the groups
plt.plot(ind, EpsilonGreedy, color="red", linewidth=2.5, linestyle="-", label="Epsilon")
plt.plot(ind, AnnealingEpsilonGreedy, color="cyan", linewidth=2.5, linestyle="-", label="Annealing Epsilon")
plt.plot(ind, Softmax, color="magenta", linewidth=2.5, linestyle="-", label="Softmax")
plt.plot(ind, AnnealingSoftmax, color="yellow", linewidth=2.5, linestyle="-", label="Annealing Softmax")
plt.plot(ind, Acc, color="pink", linewidth=2.5, linestyle="-", label="Accuracy before")
plt.legend(loc = 3)
plt.axis([0, 40, 0.7, 1])
plt.xlabel('Number of pruned Neurons')
plt.ylabel('The Accuracy')
plt.grid(True)
plt.show()
```



```
In [21]: p1 = figure(title="The Performance over the number of neurons' pruned", tools=TOOLS)

#p1.square(ind, EpsilonGreedy, legend="Epsilon Greedy", fill_color=None, line_color="red")
p1.line(ind, EpsilonGreedy, legend="Epsilon Greedy", line_color="red", line_width=2)

#p1.ellipse(ind, AnnealingEpsilonGreedy, legend="Annealing Epsilon Greedy", line_color="blue")
p1.line(ind, AnnealingEpsilonGreedy, legend="Annealing Epsilon Greedy", line_color="blue", line_width=2)

#p1.diamond(ind, Softmax, legend="Softmax", line_color="green")
p1.line(ind, Softmax, legend="Softmax", line_color="green", line_width=2)

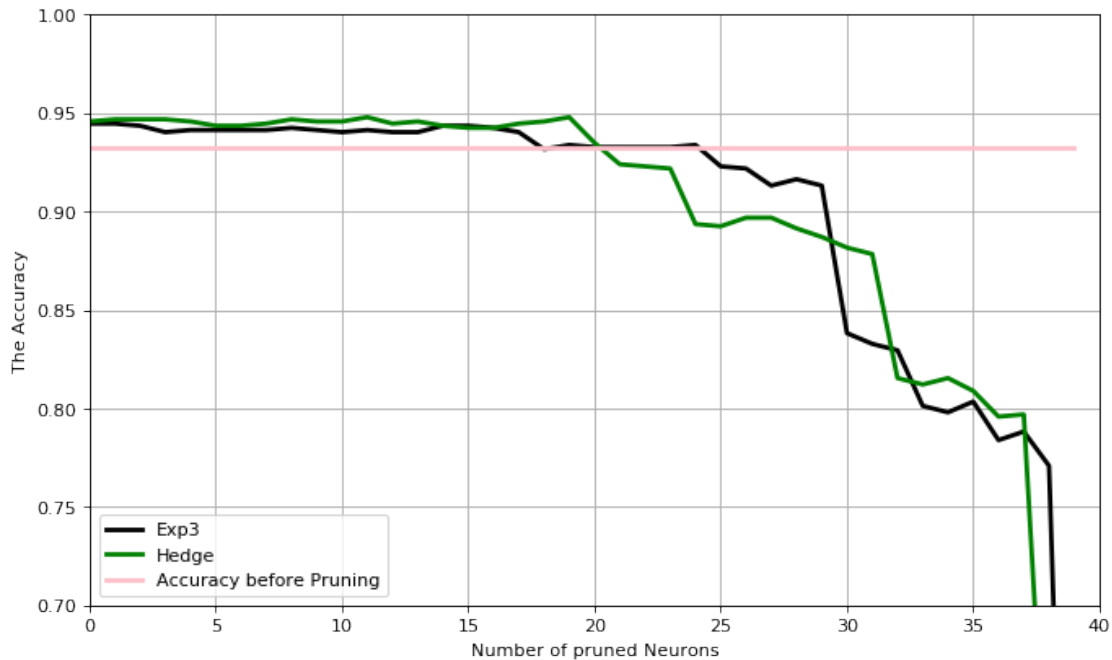
#p1.arc(ind, AnnealingSoftmax, legend="Annealing Softmax", line_color="grey", end_angle=1.57)
p1.line(ind, AnnealingSoftmax, legend="Annealing Softmax", line_color="grey", line_width=2)

p1.line(ind, Acc, legend="Accuracy", line_dash=(4, 4), line_color="orange", line_width=2)
#p1.square(ind, Hedge, legend="3*sin(x)", fill_color=None, line_color="brown")
p1.title.align = "center"

show(p1)
#show(gridplot(p1, p2, ncols=2, plot_width=400, plot_height=400)) # open a browser
```

8.4 Adversarial Bandits Hedge and EXP3

```
In [22]: fig = plt.figure(figsize=(10, 6), dpi=80)
ax = fig.add_subplot(111)
N = len(Exp3)
Acc = [Accuracy for col in range(N)]
## necessary variables
ind = np.arange(N) # the x locations for the groups
plt.plot(ind, Exp3, color="black", linewidth=2.5, linestyle="-", label="Exp3")
plt.plot(ind, Hedge, color="green", linewidth=2.5, linestyle="-", label="Hedge")
plt.plot(ind, Acc, color="pink", linewidth=2.5, linestyle="-", label="Accuracy before")
plt.legend(loc = 3)
plt.axis([0, 40, 0.7, 1])
plt.xlabel('Number of pruned Neurons')
plt.ylabel('The Accuracy')
plt.grid(True)
plt.show()
```



```
In [23]: p1 = figure(title="The Performance over the number of neurons' pruned", tools=TOOLS)

#p1.oval(ind, Exp3, legend="Exp3", line_color="black", height=0.01, width=0.01)
p1.line(ind, Exp3, legend="Exp3", line_color="black", line_width=2)

#p1.arc(ind, Hedge, legend="Hedge", line_color="yellow")
#p1.triangle(ind, Hedge, legend="Hedge", line_color="yellow")
p1.line(ind, Hedge, legend="Hedge", line_color="yellow", line_width=2)
```



```

p1.line(ind, Acc, legend="Accuracy", line_dash=(4, 4), line_color="orange", line_width=
#p1.square(ind, Hedge, legend="3*sin(x)", fill_color=None, line_color="brown")
p1.title.align = "center"

show(p1)
#show(gridplot(p1, p2, ncols=2, plot_width=400, plot_height=400)) # open a browser

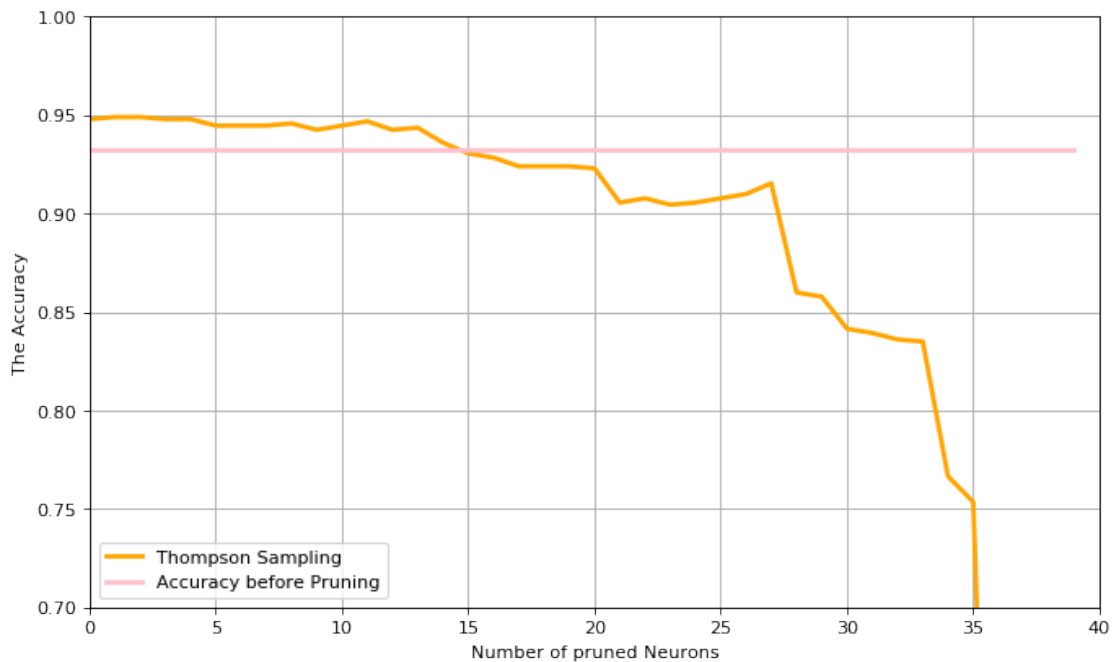
```

9 Thompson Sampling

```

In [24]: fig = plt.figure(figsize=(10, 6), dpi=80)
ax = fig.add_subplot(111)
N = len(ThompsonSampling)
Acc = [Accuracy for col in range(N)]
## necessary variables
ind = np.arange(N) # the x locations for the groups
plt.plot(ind , ThompsonSampling, color="orange", linewidth=2.5, linestyle="-", label="T")
plt.plot(ind , Acc, color="pink", linewidth=2.5, linestyle="-", label="Accuracy before")
plt.legend(loc = 3)
plt.axis([0, 40, 0.7, 1])
plt.xlabel('Number of pruned Neurons')
plt.ylabel('The Accuracy')
plt.grid(True)
plt.show()

```



```

In [25]: p1 = figure(title="The Performance over the number of neurons' pruned", tools=TOOLS)

#p1.square_cross(ind, ThompsonSampling, legend="Thompson Sampling", line_color="pink")
p1.line(ind, ThompsonSampling, legend="Thompson Sampling", line_color="pink", line_widt

p1.line(ind, Acc, legend="Accuracy", line_dash=(4, 4), line_color="orange", line_width=
#p1.square(ind, Hedge, legend="3*sin(x)", fill_color=None, line_color="brown")
p1.title.align = "center"

show(p1)
#show(gridplot(p1, p2, ncols=2, plot_width=400, plot_height=400)) # open a browser

```